


Used Car Sales Price Prediction

By C K Anand Prakash

Background

US used car market is still booming.

California used car market revenue
*Statista



Used Car Prices Are Going Up Again

Used cars are selling for higher prices at wholesale auctions as buyers return to the market, but that's driving prices up at lots.

By José Rodríguez Jr. | Published February 8, 2023 | Comments (3)




Photo: Chip Somodevilla (Getty Images)

Used car demand suddenly spiked in January, increasing wholesale values at

After a steep fall, used car prices poised to rise again

By Chris Hixson, CNN
Published 1:27 PM EST, Sat February 18, 2023

Facebook Twitter Email Print

Case Scenario

“ Hello, I am running a **used car business in California**,
I need 100 more cars to meet the market demand, and I have **three options** for a supplier.
I want to choose one who can provide me with the best cars with the **highest selling price**.
Can you help me with this? ”



Who should our client choose?

Supplier A



“I’ve statistically chosen
the 100 best cars”

Supplier B



“I sourced the best cars
through my social network”

Supplier C



“Why do you need
other suppliers, why?”

Project Overview

Objective

- Predict the sales price of cars from each supplier, and select the best one

Data set used

- Kaggle, “used-car-auction-prices”
 - Original: 558,838 rows / 16 variables
 - Cleaned: 20,000 rows / 21 variables*

Filtered California cars only

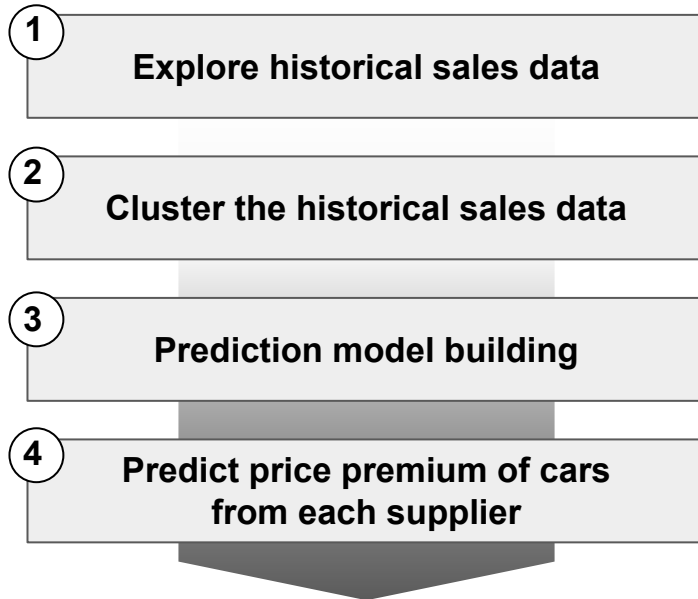
Created dependent variable: sales premium

Gathered additional variables: awarded

Methodologies applied

- Linear Regression (Lasso Regression)
- Clustering (K-means)
- Classification (Support Vector Machine)

Approach



“Why do you need this process?”

**“First, I will provide you with my sales record.
Can you show me any insights?”**



1 Explore historical sales data

2 Cluster the historical sales data

3 Prediction model building


4 Predict price premium of cars
from each supplier

Exploratory Analysis (1)


The company has sold 2,906 cars so far. Consumers seem to prefer Japanese make and below 2-year-old cars. As for individual brands, Ford was the most popular one.

Sold 2,906 cars so far...


By Fiscal Year




By Nationality



By Make



By Car Age


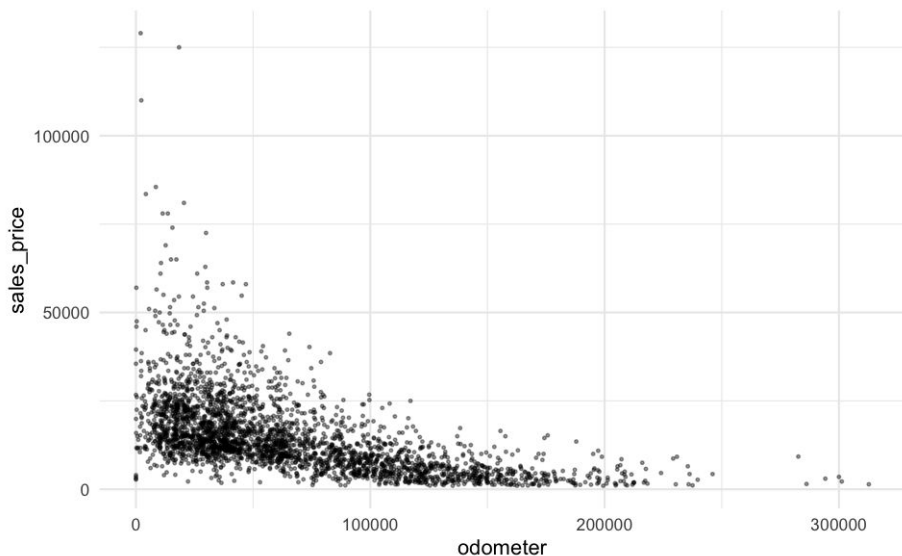


This can be viewed as the consumer preference for used cars in California

Exploratory Analysis (2)

Mileage seems to have negative impacts on the sales prices.


If a certain car is cited as a good one by the media, the car's price tends to increase.



*Award: consumer report, "best cars of the year"

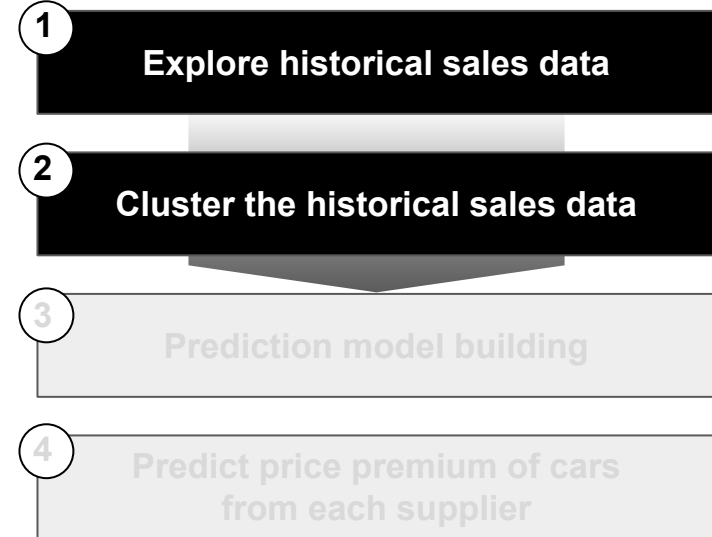
Exploratory Analysis (3)

Apart from the sales price, some cars generated higher price premiums, which means consumers were willing to pay more than the market price for those cars.



We chose the price premium as our dependent variable

**“I always felt that customers showed different buying patterns for my cars...
Can you check if my cars can be segmented?”**



Clustering


Client's cars can be clustered into two groups.

Given their heterogeneous characteristics, we decided to treat them independently.

Variables considered

```
> str(cardfl)
'data.frame': 2906 obs. of
 $ X      : int  17 18
 $ Premium : num  0.904
 $ sales_price : int  3300
 $ MMR_price : int  3650
 $ make_yr   : int  2008
 $ Age       : int  7 10
 $ sale_yr   : int  2015
 $ sale_mnt   : chr  "Jun"
 $ sale_day   : chr  "Wed"
 $ make       : chr  "Kia"
 $ national   : chr  "KOR"
 $ model      : chr  "Spec"
 $ trim       : chr  "EX"
 $ awarded    : num  0 0 0
 $ body       : chr  "Sedan"
 $ make_yr.1  : int  2008
 $ odometer   : int  93632
 $ condition  : num  2.9 2
 $ transmission: int  1 1 1
 $ ext_col    : chr  "blue"
 $ int_col    : chr  "gray"
 $ seller     : chr  "the"
```

Elbow chart



We chose 2 as the optimal size using the knee / elbow method

Segmentation result

```
> clust_data_means
```

	Group	1	Premium	sales_price	MMR_price	Age	awarded	odometer	condition	transmission
1	1	1.01	19063.24	19021.36	2.54	0.07	37733.11	3.72	0.97	
2	2	1.04	6431.09	6335.02	9.02	0.05	120572.05	2.73	0.95	

Cluster 1

Newer entry-level luxury, good condition:
Chevy Camaro, Infiniti G37



Cluster 2

Older well-driven economy, fair condition:
Honda Odyssey, Dodge Charger RT



**“I see, my cars can be clustered into two groups.
Now let’s build a prediction model”**



Prediction Model Building

We built two different Lasso regression models for each cluster.

Issues

1. Information discrepancy

- Our data has sales related Information, such as sales price and sales data.
But dealer's data set doesn't

👉 **Remove sales-related variables on modeling**

2. Still too many variables

- 103 variables in total, after converting all categorical variables to dummy variables (high risk of multicollinearity)

👉 **Take alternative approach**

Alternative Approach

1. Feature selection with the Client's domain knowledge

- "MMR price, Age, Odometer, Condition, Awarded, and Transmission is everything"
- R^2 : 0.35, RMSE: 0.09738 (for cluster 1)

2. Lasso regression

- R^2 : 0.61, RMSE: 0.0979 (for cluster 1)

👉 **Lasso regression showed better model fit and prediction power**

Model building

Model for Cluster 1

```
> coef(lasso.c1)
103 x 1 sparse Matrix of class "dgCMatrix"
               s0
(Intercept)    6.283164e-01
MMR_price      -9.798730e-06
Age            7.374449e-03
awarded        .
odometer      -4.037326e-08
condition      1.444955e-01
transmission    .
makeAcura      .
```

Model for Cluster 2

```
> coef(lasso.c2)
103 x 1 sparse Matrix of class "dgCMatrix"
               s0
(Intercept)    6.344006e-01
MMR_price      -8.707714e-06
Age            7.032860e-03
awarded        .
odometer        .
condition      1.396396e-01
transmission    .
makeAcura      .
makeAston.Martin .
makeAudi        .
```

Prediction Model Building



“Why do you need
two different models? why?”

		1. Single model for the entire data set	2. Two models optimized for each cluster
Client's domain knowledge model	R-Squared	0.2226	Cluster1: 0.3545 Avg Cluster2: 0.2179 0.286
	RMSE	0.1764	Cluster1: 0.0995 Avg Cluster2: 0.2468 0.173
Lasso Regression	R-Squared	0.6531	Cluster1: 0.6279 Avg Cluster2: 0.756 0.691
	RMSE	0.1745	Cluster1: 0.0983 Avg Cluster2: 0.2424 0.170

Two model approach
shows better prediction
power and model fit

**“Now we have prediction models.
Tell me which supplier should I choose?”**



**“This is a tense moment.
Do you need a joke?”**

Prediction (1)

First, we classified each supplier's cars and applied the cluster-optimized prediction model. And the winner is...

Support Vector Machine with “radial” kernel

Modeled with the top 10 variables highly correlated with the cluster variable




Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	907	31
1	7	498

Accuracy : 0.9737
95% CI : (0.964, 0.9813)
No Information Rate : 0.6334
P-Value [Acc > NIR] : < 2.2e-16

We chose SVM since the dataset is high-dimensional, and the client favored robustness of a model over interpretability

Classify each supplier's cars



Cluster 1
Cluster 2

“My social network said I will get the deal”

Prediction (2) - Conclusion

Supplier A shows the highest and most stable price premium.
Therefore, we highly recommend him for your business.

Forecast of price premium of 100 cars

Mean : 1.023
Sd : 0.095



Mean : 1.034
Sd : 0.069




Mean : 1.014
Sd : 0.079



"I will quit and
teach marketing
instead"




Distribution of price premium



**Difference is marginal
since, originally, they are from the same data set**

Lessons Learned



“ Dependent variables are something that can be created ”

- Sometimes playing around with variables can provide you with different insights

“ Every population has heterogeneous nature ”

- It is the era of “micro-segmentation” modeling on segmented groups can be helpful

“ It’s a message, not a model ”

- In real-world business settings, decision makers tend to focus more on messages than models

Appendix

[I]	Data set and additional exploratory analysis	19-21
[II]	Clustering	22
[III]	Prediction Modeling	23-25
[IV]	Classification	26-27

Data Set

Original Data Set


Kaggle, “used-car-auction-prices” Original: 558,838 rows / 16 variables

year	make	model	trim	body	transmission	vin	state	condition	odometer	color	interior	seller	mmr	sellingprice	saledate
2015	Kia	Sorento	LX	SUV	automatic	5xyktca69fg566472	ca	5	16639	white	black	kia motors	20500	21500	Tue Dec 16 2014 13:00
2015	Kia	Sorento	LX	SUV	automatic	5xyktca69fg561319	ca	5	9393	white	beige	kia motors	20800	21500	Tue Dec 16 2014 13:00
2014	BMW	3 Series	328i SULE	Sedan	automatic	wba3c1c51ek116351	ca	4.5	1331	gray	black	financial services	31900	30000	Thu Jan 23 2015 12:00
2015	Volvo	S60	T5	Sedan	automatic	yv1612tb4f1310987	ca	4.1	14282	white	black	volvo na ri	27500	27750	Thu Jan 23 2015 12:00
2014	BMW	6 Series G	650i	Sedan	automatic	wba6b2c57ed129731	ca	4.3	2641	gray	black	financial services	66000	67000	Thu Dec 18 2014 12:00
2015	Nissan	Altima	2.5 S	Sedan	automatic	1n4al3ap1fn326013	ca	1	5554	gray	black	enterprise	15350	10900	Tue Dec 16 2014 13:00
2014	BMW	M5	Base	Sedan	automatic	wbsfv9c51ed593089	ca	3.4	14943	black	black	the hertz co	69000	65000	Wed Dec 17 2014 12:00
2014	Chevrolet	Cruze	1LT	Sedan	automatic	1g1pc5sb2e7128460	ca	2	28617	black	black	enterprise	11900	9800	Tue Dec 16 2014 13:00
2014	Audi	A4	2.0T Premium	Sedan	automatic	wauffaf13en030343	ca	4.2	9557	white	black	audi mississauga	32100	32250	Thu Dec 18 2014 12:00
2014	Chevrolet	Camaro	LT	Convertible	automatic	2g1fb3d37e9218789	ca	3	4809	red	black	d/m auto	26300	17500	Tue Jan 20 2015 12:00
2014	Audi	A6	3.0T Prestige	Sedan	automatic	wauhgafc0en062916	ca	4.8	14414	black	black	desert auto	47300	49750	Tue Dec 16 2014 13:00
2015	Kia	Optima	LX	Sedan	automatic	5xxgm4a73fg353538	ca	4.8	2034	red	tan	kia motors	15150	17700	Tue Dec 16 2014 13:00
2015	Ford	Fusion	SE	Sedan	automatic	3fa6p0hdxr145753	ca	2	5559	white	beige	enterprise	15350	12000	Tue Jan 20 2015 12:00
2015	Kia	Sorento	LX	SUV	automatic	5xyktca66fg561407	ca	5	14634	silver	black	kia motors	20600	21500	Tue Dec 16 2014 13:00
2014	Chevrolet	Cruze	2LT	Sedan	automatic	1g1pe5sbxe7120097	ca		15686	blue	black	avis rac/sa	13900	10600	Tue Dec 16 2014 13:00
2015	Nissan	Altima	2.5 S	Sedan	automatic	1n4al3ap5fc124223	ca	2	11398	black	black	enterprise	14750	14100	Tue Dec 16 2014 13:00
2015	Hyundai	Sonata	SE	Sedan		5npe24af4fh001562	ca		8311	red	???"avis tra"	15200	4200	Tue Dec 16 2014 13:00	
2014	Audi	Q5	2.0T Premium	SUV	automatic	wa1ffa1fxea085074	ca	4.9	7983	white	black	audi north	37100	40000	Thu Dec 18 2014 12:00
2014	Chevrolet	Camaro	LS	Coupe	automatic	2g1fa1e39e9134494	ca	1.7	13441	black	black	wells fargo	17750	17000	Tue Dec 16 2014 13:00
2014	BMW	6 Series	650i	Convertible	automatic	wbays9c53ed169260	ca	3.4	8819	black	black	the hertz co	68000	67200	Wed Dec 17 2014 12:00
2015	Chevrolet	Impala	LTZ	Sedan	automatic	2g1165s30f9103921	ca	1.9	14538	silver	black	enterprise	24300	7200	Tue Jul 14 2015 12:00
2014	BMW	5 Series	528i	Sedan	automatic	wba5a5c51ed501631	ca	2.9	25969	black	black	financial services	34200	30000	Tue Feb 10 2015 12:00
2014	Chevrolet	Camaro	LT	Convertible	automatic	2n1fb3d31e9134662	ca		33450	black	black	avis rac/sa	20100	14700	Tue Dec 16 2014 13:00

1. Delete NA values
2. Delete invalid values
 - e.g., sales_price = white
3. Delete outliers
 - e.g. price = 30,000,000
4. Create new variables
 - Premium (sales_price / MMR_price)
 - Awarded (consumer reports)
 - National (nationality of makes)
5. Dummy variable conversion
 - e.g., transmission, state
6. Sampled 20,000 observations
7. Filtered California cars only


Sales price vs. Condition

“Condition” variable is an important indicator of sales price in the used car market.
Used cars under greater condition is more likely to be sold in higher prices.



Correlation between features

- “Odometer” is **negatively** related to sales prices in the used car market.
- “Mnf price” is **positively** related to sales prices.
- “Make_yr” and “condition” variables are also an important indicator of sales prices.



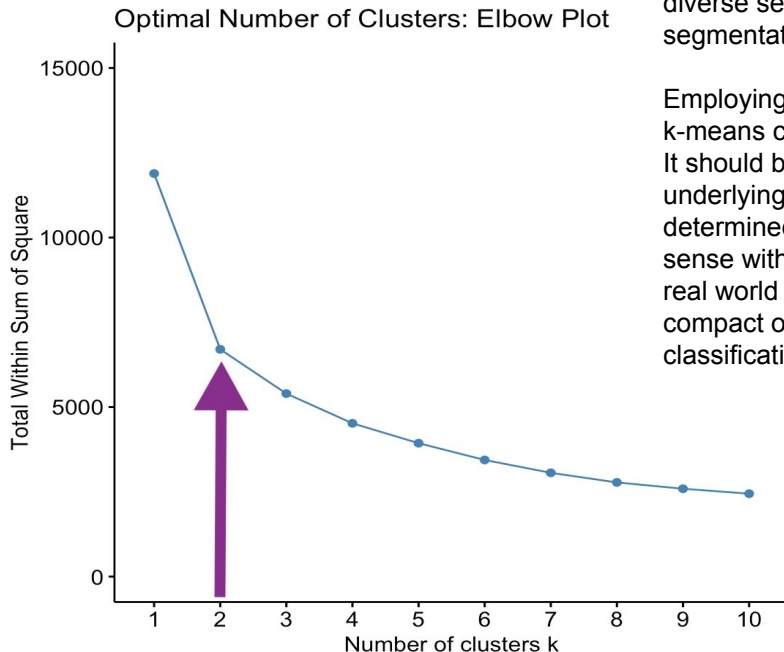
Clustering

#Elbow plot

```
fviz_nbclust(x=clusterdf3.1std, FUNcluster = kmeans, nstart=100, method="wss", k.max = 10) +  
  labs(title="Optimal Number of Clusters: Elbow Plot") +  
  coord_cartesian(ylim=c(0,15000)) + geom_line(size=2) # --> looks like K=2 optimal
```

As shown in the exploratory analysis of this project, the cars in this dataset represent a diverse set of vehicular demographics. As such, it was necessary to examine the natural segmentation of the dataset.

Employing the knee/elbow approach to determine the optimal number of clusters k for k -means clustering, we selected 2 for k , as the elbow/knee is distinctly accentuated at $k=2$. It should be noted that we did also attempt higher values of k and did examine the underlying averages for the segments created for these higher k values, but we ultimately determined that two segments represented a clear delineation across groups that made sense within the scope of this project. Applying the decision making within clustering to our real world setting, both the elbow/knee method as well as the benefit of having more compact operational complexity (from two clusters, as opposed to three or four per classification model) pointed towards using 2 clusters.



#	Group.1	Premium	sales_price	MMR_price	Age	awarded	odometer	condition	transmission
# 1	1	1.01	19063.24	19021.36	2.54	0.07	37733.11	3.72	0.97
# 2	2	1.04	6431.09	6335.02	9.02	0.05	120572.05	2.73	0.95

#	Group.1	Premium	sales_price	MMR_price	Age	awarded	odometer	condition	transmission
# 1	1	0.98	12625.51	12983.74	3.33	0.05	51798.69	3.06	0.97
# 2	2	1.06	5834.92	5635.82	9.69	0.05	129316.98	2.72	0.95
# 3	3	1.02	24044.02	23667.11	2.25	0.09	29989.77	4.23	0.98

Prediction Modeling (1)

Client's Model Summary for train data for cluster 1

```
> summary(fit0)
```

Call:

```
lm(formula = Premium ~ MMR_price + Age + awarded + odometer +  
    condition + transmission, data = train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.67440	-0.04945	0.00075	0.05276	0.57193

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.435e-01	2.232e-02	28.823	< 2e-16 ***
MMR_price	-1.479e-06	2.673e-07	-5.534	3.76e-08 ***
Age	1.634e-02	2.176e-03	7.512	1.07e-13 ***
awarded	9.647e-03	1.079e-02	0.895	0.371206
odometer	5.580e-07	1.612e-07	3.461	0.000554 ***
condition	8.503e-02	3.542e-03	24.006	< 2e-16 ***
transmission	1.120e-02	1.655e-02	0.677	0.498542

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1008 on 1320 degrees of freedom
Multiple R-squared: 0.3512, Adjusted R-squared: 0.3483
F-statistic: 119.1 on 6 and 1320 DF, p-value: < 2.2e-16

Client's Model Summary for entire data for cluster 1

```
> summary(fit.ceo1)
```

Call:

```
lm(formula = Premium ~ MMR_price + Age + awarded + odometer +  
    condition + transmission, data = a1)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.67723	-0.04817	0.00085	0.05342	0.57383

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.569e-01	1.858e-02	35.365	< 2e-16 ***
MMR_price	-1.524e-06	2.303e-07	-6.620	4.66e-11 ***
Age	1.745e-02	1.785e-03	9.779	< 2e-16 ***
awarded	1.268e-02	8.924e-03	1.421	0.155493
odometer	4.614e-07	1.314e-07	3.511	0.000457 ***
condition	8.258e-02	2.903e-03	28.449	< 2e-16 ***
transmission	7.912e-03	1.425e-02	0.555	0.578735

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.09977 on 1889 degrees of freedom
Multiple R-squared: 0.3545, Adjusted R-squared: 0.3524
F-statistic: 172.9 on 6 and 1889 DF, p-value: < 2.2e-16

Prediction Modeling (2)

Client's Model Summary for train data for cluster 2

```
> summary(fit0)
```

```
Call:
lm(formula = Premium ~ MMR_price + Age + awarded + odometer +
    condition + transmission, data = train)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-0.81445 -0.14524 -0.00665  0.12935  1.78719
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	5.943e-01	7.627e-02	7.792	2.40e-14	***
MMR_price	-1.403e-05	3.027e-06	-4.633	4.29e-06	***
Age	1.190e-02	3.951e-03	3.011	0.0027	**
awarded	-7.402e-03	4.453e-02	-0.166	0.8680	
odometer	-2.236e-07	1.601e-07	-1.396	0.1630	
condition	1.665e-01	1.376e-02	12.103	< 2e-16	***
transmission	4.747e-03	4.306e-02	0.110	0.9122	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2548 on 700 degrees of freedom
Multiple R-squared: 0.2212, Adjusted R-squared: 0.2145
F-statistic: 33.13 on 6 and 700 DF, p-value: < 2.2e-16

Client's Model Summary for entire data for cluster 2

```
> summary(fit.ceo2)
```

```
Call:
lm(formula = Premium ~ MMR_price + Age + awarded + odometer +
    condition + transmission, data = a2)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-0.81175 -0.15047 -0.00539  0.13223  1.79630
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	5.998e-01	6.690e-02	8.965	< 2e-16	***
MMR_price	-1.297e-05	2.591e-06	-5.008	6.49e-07	***
Age	1.305e-02	3.311e-03	3.942	8.65e-05	***
awarded	2.593e-02	3.905e-02	0.664	0.5069	
odometer	-2.885e-07	1.499e-07	-1.925	0.0545	.
condition	1.606e-01	1.152e-02	13.940	< 2e-16	***
transmission	3.835e-03	3.868e-02	0.099	0.9211	


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2578 on 1003 degrees of freedom
Multiple R-squared: 0.2079, Adjusted R-squared: 0.2031
F-statistic: 43.87 on 6 and 1003 DF, p-value: < 2.2e-16

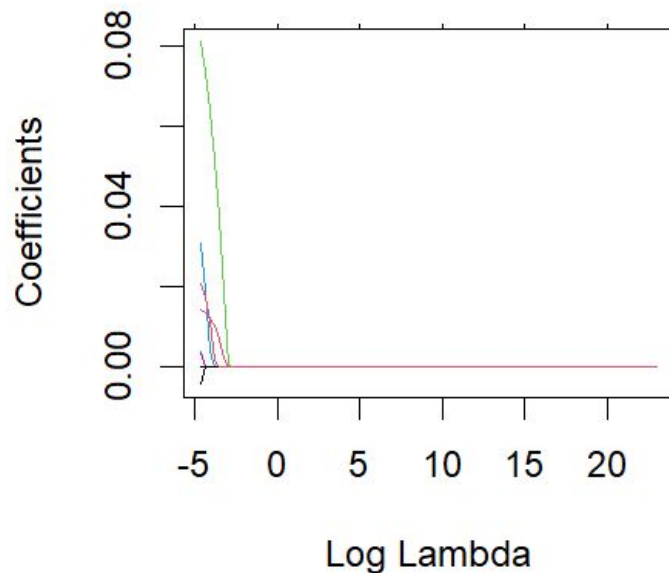
Prediction Modeling (3)

Lasso Regression

To find the best lambda, we first tried a loop method, and then conducted the cross validation



Using the best lambda, we built the lasso regression model, the number of variables included in the model reduced to 11.



Classifications(1)

Random Forest

```
install.packages('randomForest')
library(randomForest)
set.seed(223344)
bag.train.100 <- randomForest(ClusterNo ~ .,
                             data = r2,
                             mtry = 12, ntree = 100,
                             importance = TRUE)

bag.train.100
# Mean of squared residuals: 0.0235618
# % Var explained: 89.61
```

```
bag.train.50 <- randomForest(ClusterNo ~ .,
                             data = r2,
                             mtry = 12, ntree = 50,
                             importance = TRUE)

bag.train.50
# Mean of squared residuals: 0.02322343
# % Var explained: 89.76
```

The accuracy of random forest is about 89.76%,
with 50 number of trees, MSE=0.0232

```
Call:
randomForest(formula = ClusterNo ~ ., data = r2, mtry = 12, ntree = 50, importance = TRUE)
Type of random forest: regression
Number of trees: 50
No. of variables tried at each split: 12

Mean of squared residuals: 0.02322343
% Var explained: 89.76
```

Logistic Regression Confusion Matrix

```
> yhat.test.class <- ifelse(yhat.test > 0.5, 1, 0)
> table(r2[test, ]$ClusterNo,
+       yhat.test.class,
+       dnn = c("Actual", "Predicted"))
```

	Predicted	
Actual	0	1
0	908	32
1	31	482

This confusion matrix shows the performance of the classification model on the test set. The rows represent the actual class labels, while the columns represent the predicted class labels.

Looking at the table, we can see that the model predicted 482 instances as positive (1) when they were actually positive, and 908 instances as negative (0) when they were actually negative. However, the model incorrectly predicted 31 instances as negative when they were actually positive (false negatives) and 32 instances as positive when they were actually negative (false positives).

$\text{accuracy} = (908 + 482) / 1453 = 0.9566$

The accuracy is quite high, which means that the model is able to correctly classify most instances in the dataset.

Classifications(2)

SVM

```
library(e1071)

set.seed(100)

dat.train$ClusterNo <- factor(dat.train$ClusterNo)

# Find the best svm model

tune.out <- tune(svm, ClusterNo ~ Age + odometer + MMR_price + condition + int_colblack + int_colgray + nationalKOR + bodySedan + makeInfiniti + bodyCab,
  data = dat.train, kernel = "radial",
  ranges = list(cost = c(0.01, 0.1, 1, 10, 100, 1000), gamma = c(0.5, 1, 2, 3, 4)))

# This is the best SVM model

svm.best <- tune.out$best.model

svm.pred <- predict(svm.best, dat.test, type="class")

table(actual=dat.test[[1]], predict=svm.pred)

library(caret)

confusionMatrix(factor(dat.test[[1]]), svm.pred)
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	907	31
1	7	498


Accuracy : 0.9737
95% CI : (0.964, 0.9813)
No Information Rate : 0.6334
P-Value [Acc > NIR] : < 2.2e-16

- We select the top 10 highest correlated variables, which are: Age, odometer, MMR_price, condition, int_colblack, int_colgray, nationalKOR, bodySedan, makeInfiniti, bodyCab
- SVM model has a high accuracy of 97.37% on the test data, which means that the model was able to correctly predict the target variable for 97.37% of the observations in the test data set
- Therefore, we choose the SVM model because it has the highest accuracy of the three classification models.

Prediction


Data preparation

We randomly assigned 100 observations to each supplier



Classification

And then classified each supplier's cars into two groups



Forecasting

Finally, we applied the prediction model to each cluster, and get the results

Prediction for the price premium

	Mean	Standard Deviation
Supplier A	1.034	0.069
Supplier B	1.023	0.095
Supplier C	1.014	0.079